

# Step-by-Step Implementation Plan: Process vs. Outcome Supervision for Trustworthy and Interpretable LLM Reasoning

## Project Overview

**Goal:** Compare process-based supervision (step-by-step reasoning) vs. outcome-based supervision (final answers only) in terms of (a) factual reliability and (b) interpretability/auditability for multi-step reasoning tasks.

**Timeline:** 6 weeks

**Scope:** Evaluation study (no model training required)

**Target:** 150-200 math problems, 50-70 annotated in depth

---

## Phase 1: Setup and Data Collection (Week 1-2)

### Week 1: Environment Setup and Model Selection

#### Step 1.1: Set Up Development Environment

##### Tasks:

- Set up Python environment (Python 3.9+)
- Install core libraries:

```
bash
```

```
pip install transformers torch datasets openai anthropic pandas numpy scikit-learn jupyter matplotlib seaborn --break-system
```

- Set up API access:
  - Option A: OpenAI API (GPT-4o-mini for cost-efficiency)
  - Option B: Anthropic API (Claude 3.5 Sonnet)
  - Option C: Open-weight model via HuggingFace (Qwen2.5-Math-7B-Instruct or Llama-3.1-8B-Instruct)
- Create project directory structure:

```
project/
├── data/
│   ├── raw/
│   ├── processed/
│   └── annotations/
├── outputs/
├── scripts/
├── notebooks/
└── results/
```

## Resources:

- HuggingFace Transformers documentation
- API documentation (OpenAI/Anthropic)

**Deliverable:** Working development environment with API access confirmed

---

## Step 1.2: Select and Download Dataset

### Tasks:

Download GSM8K dataset from HuggingFace:

```
python
```

```
from datasets import load_dataset
gsm8k = load_dataset("openai/gsm8k", "main")
```

Sample 150-200 problems from test set:

- Use stratified sampling if possible (by difficulty/length)
- Ensure problems have clear ground-truth answers
- Verify solutions have step-by-step reasoning

Create data manifest (CSV with problem\_id, question, answer, solution\_steps)

### Alternative datasets if GSM8K insufficient:

- MATH dataset (more challenging): `load_dataset("hendrycks/competition_math")`
- SVAMP (variations): `load_dataset("ChilleD/SVAMP")`

## Resources:

- GSM8K: <https://huggingface.co/datasets/openai/gsm8k>
- MATH: [https://huggingface.co/datasets/hendrycks/competition\\_math](https://huggingface.co/datasets/hendrycks/competition_math)

**Deliverable:** CSV file with 150-200 selected problems including ground truth

---

### Step 1.3: Design Prompts for Both Conditions

#### Tasks:

##### Create **Outcome-only prompt template**:

Answer the following math problem. Provide ONLY the final numerical answer, without showing any work or reasoning.

Problem: {question}

Answer:

##### Create **Process prompt template** (standard CoT):

Solve the following math problem step by step. Show all your reasoning clearly, then provide the final answer.

Problem: {question}

Solution:

##### Create **Structured Process prompt template** (Process+):

Solve the following math problem using the following format:

Step 1: [First reasoning step]

Step 2: [Second reasoning step]

...

Final Answer: [Numerical answer]

Problem: {question}

Solution:

##### Test prompts on 5 sample problems to verify format

#### Resources:

- Chain-of-Thought Prompting paper examples
- OpenAI/Anthropic API prompt engineering guides

**Deliverable:** Three validated prompt templates

---

## Week 2: Data Generation

### Step 2.1: Generate Model Responses

#### Tasks:

Write data collection script with:

- API rate limiting and retry logic
- Progress tracking and checkpointing
- Response parsing and storage
- Error handling

Generate responses for ALL conditions:

```
python

for problem in problems:
    # Outcome condition
    outcome_response = query_model(outcome_prompt.format(question=problem))

    # Process condition
    process_response = query_model(process_prompt.format(question=problem))

    # Process+ condition (optional)
    structured_response = query_model(structured_prompt.format(question=problem))

    # Store with problem_id, condition, response, timestamp
```

Implement response storage format:

```
python

{
    "problem_id": "gsm8k_001",
    "question": "...",
    "ground_truth": "42",
    "ground_truth_solution": "...",
    "outcome_response": "42",
    "process_response": "Step 1: ... Step 2: ... Therefore, 42",
    "structured_response": "Step 1: ... Final Answer: 42",
    "timestamp": "2025-01-15T10:30:00"
}
```

**Cost estimation (OpenAI GPT-4o-mini):**

- $\sim 200$  problems  $\times 3$  conditions  $\times \sim 500$  tokens avg =  $\sim 300K$  tokens
- Input:  $\sim \$0.45$ , Output:  $\sim \$1.80$ , Total:  $\sim \$2.25$

## Resources:

- OpenAI API: <https://platform.openai.com/docs/api-reference>
- Anthropic API: <https://docs.anthropic.com/claude/reference>

**Deliverable:** JSON/CSV file with all model responses for all conditions

---

## Step 2.2: Parse and Extract Final Answers

### Tasks:

- Implement answer extraction logic:
  - For outcome condition: direct extraction
  - For process/structured: parse last line or "Final Answer:" marker
  - Handle edge cases (no answer, multiple numbers, formatting issues)
- Create answer normalization function:

```
python
```

```
def normalize_answer(text):
    # Extract numerical value
    # Handle units, fractions, decimals
    # Standardize format
    return normalized_value
```

- Validate extraction quality manually on sample (20 problems)
- Create processed dataset with extracted answers:

```
python
```

```
{  
    "problem_id": "gsm8k_001",  
    "ground_truth_normalized": 42.0,  
    "outcome_answer": 42.0,  
    "process_answer": 42.0,  
    "structured_answer": 42.0,  
    "outcome_correct": True,  
    "process_correct": True,  
    "structured_correct": True  
}
```

## Resources:

- Regex patterns for number extraction
- GSM8K evaluation code: <https://github.com/openai/grade-school-math>

**Deliverable:** Processed dataset with extracted and normalized answers

---

## Phase 2: Metrics Implementation (Week 3)

### Week 3: Implement Core Metrics

#### Step 3.1: Calculate Final Answer Accuracy

##### Tasks:

- Implement accuracy calculation:

```
python  
  
def calculate_accuracy(predictions, ground_truth):  
    correct = sum([pred == gt for pred, gt in zip(predictions, ground_truth)])  
    return correct / len(predictions)
```

- Compute accuracy for each condition:

- Accuracy(Outcome)
- Accuracy(Process)
- Accuracy(Process+)

- Perform statistical significance testing:

```
python
```

```

from scipy.stats import mcnemar
# McNemar's test for paired binary outcomes
contingency_table = [[both_correct, process_correct_outcome_wrong],
                     [outcome_correct_process_wrong, both_wrong]]
statistic, pvalue = mcnemar(contingency_table)

```

- Create accuracy comparison visualization

**Deliverable:** Accuracy metrics with statistical significance tests

---

### Step 3.2: Error-Type Analysis (Subset)

**Tasks:**

- Select 30-50 incorrectly answered problems across conditions
- Define error taxonomy:
  1. **Arithmetic error:** Calculation mistake (e.g.,  $3 \times 5 = 11$ )
  2. **Conceptual error:** Wrong approach or formula
  3. **Reading comprehension error:** Misunderstood problem
  4. **Incomplete reasoning:** Missing steps or logic gaps
  5. **Formatting/parsing error:** Answer present but not extracted
- Manually label errors by type for each condition
- Create error distribution comparison:

```

python

import matplotlib.pyplot as plt
import seaborn as sns

# Create stacked bar chart showing error types per condition

```

**Deliverable:** Error taxonomy with distribution across conditions

---

### Step 3.3: Hallucination Rate in Process Responses

**Tasks:**

- Select 50 process/structured responses for hallucination analysis
- Define hallucination categories:
  1. **Factual error:** Objectively false statement ( $3 \times 5 = 11$ )

**2. Irrelevant information:** Introduces unrelated facts

**3. Logical inconsistency:** Contradicts earlier steps

**4. Confabulation:** Makes up non-existent details

Binary label: Does response contain  $\geq 1$  hallucination?

Calculate hallucination rate:

```
python
```

```
hallucination_rate = num_with_hallucinations / total_responses
```

Correlate with correctness:

- Hallucination rate among correct answers
- Hallucination rate among incorrect answers

**Deliverable:** Hallucination rate metric with analysis

---

## Phase 3: Interpretability Metrics (Week 4)

### Week 4: Deep Annotation for Interpretability

#### Step 4.1: Sample Selection for Deep Annotation

**Tasks:**

Select 50-70 problems for intensive annotation:

- Stratified by correctness (correct/incorrect answers)
- Stratified by condition (ensure coverage)
- Include diverse problem types

Prepare annotation interface (spreadsheet or custom tool):

- Problem text
- Ground truth solution (step-by-step)
- Model response (with steps highlighted)
- Annotation fields (see below)

**Deliverable:** Annotated problem set (50-70 problems)

---

#### Step 4.2: Implement Step Correctness Score

**Tasks:**

Segment process responses into individual steps:

- For structured: parse by "Step N:" markers
- For unstructured: manual or heuristic segmentation (by sentence/line)

Design annotation rubric:

Score 0: Incorrect or unjustified

- Factual error (wrong calculation, false premise)
- Missing justification
- Illogical inference

Score 1: Partially correct or incomplete

- Correct direction but imprecise
- Partially justified
- Minor arithmetic errors with correct approach

Score 2: Correct and well-justified

- Factually accurate
- Clear reasoning
- Properly follows from previous steps

Annotate each step for 50-70 problems

Calculate Step Correctness Score:

```
python
```

```
step_correctness_score = mean([step_scores for all problems])
# Average score per problem, then average across problems
```

Inter-annotator reliability check:

- If 2 annotators: Calculate Cohen's kappa on 20% overlap
- Resolve disagreements through discussion

**Deliverable:** Step correctness scores with inter-rater reliability

---

### Step 4.3: Calculate Faithfulness / Expert Alignment Score

**Tasks:**

For each problem, align model steps with expert solution steps:

Expert Step 1: Find total items →  $3 + 5 = 8$

Model Step 1: Add the quantities →  $3 + 5 = 8$

Alignment: Match (Score 2)

Expert Step 2: Divide by cost →  $8 / 2 = 4$

Model Step 2: Multiply by cost →  $8 \times 2 = 16$

Alignment: Different operation (Score 0)

Design faithfulness rubric:

Score 0: Different and incorrect/irrelevant

- Model uses different operation and gets wrong result
- Introduces irrelevant reasoning

Score 1: Similar operation but imprecise/partially wrong

- Correct operation type but sloppy execution
- Right general approach but inefficient

Score 2: Essentially same operation and correct

- Model reasoning matches expert reasoning
- May use different words but same logic

Calculate Faithfulness Score:

```
python
```

```
faithfulness_score = mean([alignment_scores for all problems])
```

Optional: Compute automatic similarity proxy using embeddings:

```
python
```

```
from sentence_transformers import SentenceTransformer
model = SentenceTransformer('all-MiniLM-L6-v2')

expert_embedding = model.encode(expert_step)
model_embedding = model.encode(model_step)
similarity = cosine_similarity([expert_embedding], [model_embedding])[0][0]
```

**Deliverable:** Faithfulness/expert alignment scores

---

#### Step 4.4: Human Auditability Evaluation

**Tasks:**

- Design auditability questionnaire (5-point Likert scale): **Clarity:** 1 - Very unclear / 2 - Unclear / 3 - Neutral / 4 - Clear / 5 - Very clear "How easy is it to understand what the model is doing at each step?" **Verification Effort:** 1 - Very difficult to verify / 5 - Very easy to verify "How much effort would it take to check if each step is correct?" **Coherence:** 1 - Incoherent / 5 - Highly coherent "Do the steps flow logically from one to the next?"
- Conduct evaluation with 2-3 raters (including yourself):
  - Rate same 50-70 problems
  - Provide brief justification for scores
- Calculate inter-rater reliability (Fleiss' kappa or ICC)
- Aggregate scores:

```
python
```

```
mean_clarity = mean([clarity_scores across raters and problems])
mean_verification_effort = mean([effort_scores across raters and problems])
mean_coherence = mean([coherence_scores across raters and problems])
```

- Compare across conditions (Outcome has no reasoning to audit)

**Deliverable:** Human auditability scores with reliability metrics

---

## Phase 4: Analysis and Qualitative Study (Week 5)

### Week 5: Comprehensive Analysis

#### Step 5.1: Statistical Analysis

**Tasks:**

- Compare metrics across conditions:

```
python
```

```

import pandas as pd
from scipy import stats

# Create comparison DataFrame
results_df = pd.DataFrame({
    'Condition': ['Outcome', 'Process', 'Process+'],
    'Accuracy': [outcome_acc, process_acc, structured_acc],
    'Step_Correctness': [None, process_step_corr, structured_step_corr],
    'Faithfulness': [None, process_faith, structured_faith],
    'Clarity': [None, process_clarity, structured_clarity],
    'Verification_Effort': [None, process_verify, structured_verify]
})

# Statistical tests
# Paired t-test for accuracy
t_stat, p_value = stats.ttest_rel(outcome_correct, process_correct)

# Effect sizes (Cohen's d)
cohens_d = (mean_process - mean_outcome) / pooled_std

```

- Create correlation analysis:
  - Accuracy vs. Step Correctness
  - Accuracy vs. Faithfulness
  - Step Correctness vs. Faithfulness
  - Verification Effort vs. Accuracy
- Generate summary statistics table

**Deliverable:** Statistical analysis report with significance tests

---

## Step 5.2: Qualitative Failure Case Analysis

### Tasks:

- Identify interesting failure patterns:
  1. **Correct answer, garbage reasoning:** Final answer correct but steps invalid
  2. **Wrong answer, good reasoning:** Sound approach but arithmetic error
  3. **Hallucination with confidence:** Confidently stated false facts
  4. **Incomplete reasoning:** Jumps steps, still reaches correct answer
- Select 10-15 representative examples for each pattern

- Document detailed case studies:

Case ID: gsm8k\_042

Problem: [text]

Ground Truth: 15

Outcome Response: "15"

Process Response: "First, we calculate  $3 \times 5 = 15$ . Then..."

Analysis:

- Final answer correct in both conditions
- Process condition shows arithmetic error in Step 1
- Despite error, arrives at correct answer (compensating error in Step 3)
- Demonstrates process supervision reveals hidden errors

- Create categorized failure taxonomy with examples

**Deliverable:** Qualitative analysis document with case studies

---

### Step 5.3: Visualization Creation

**Tasks:**

- Create comprehensive visualizations:

- 1. Accuracy comparison bar chart**
- 2. Error type distribution (stacked bars)**
- 3. Step correctness distribution (box plots)**
- 4. Faithfulness vs. Accuracy scatter plot**
- 5. Verification effort comparison**
- 6. Correlation heatmap (all metrics)**

- Use consistent color scheme and style:

python

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style("whitegrid")
colors = {'Outcome': '#FF6B6B', 'Process': '#4ECDC4', 'Process+': '#95E1D3'}
```

- Add error bars and statistical annotations where relevant

Export high-resolution figures (300 DPI for paper)

**Deliverable:** Complete set of publication-quality visualizations

---

## Phase 5: Synthesis and Documentation (Week 6)

### Week 6: Final Report and Presentation

#### Step 6.1: Connect Findings to Legal/Regulatory Context

##### Tasks:

Analyze results through regulatory lens: **If Process Supervision Shows Higher Reliability:**

- Discuss implications for GDPR Article 22 (right to explanation)
- Connect to AI Act transparency requirements
- Argue for process-based explanations in high-stakes domains

**If Process Supervision Shows Higher Interpretability:**

- Discuss auditability for algorithmic accountability
- Connect to requirements for meaningful human oversight
- Argue transparency enables non-arbitrary decision-making

**If Reasoning Traces Are Unfaithful:**

- Warn about liability risks of misleading explanations
- Discuss dangers of "illusion of explainability"
- Recommend caution in using CoT for compliance

Reference specific legal frameworks:

- EU GDPR (Articles 13-15, 22)
- EU AI Act (Articles 13, 52)
- U.S. Administrative Procedure Act (§ 706)
- Emerging case law on algorithmic accountability

Draft 2-3 page legal implications section

**Deliverable:** Legal/regulatory implications analysis

---

#### Step 6.2: Write Final Report

##### Tasks:

Structure report following standard format: **1. Introduction (2 pages)**

- Research question and significance

- Connection to Responsible AI and Law

## **2. Related Work (2-3 pages)**

- Process supervision literature
- Interpretability and faithfulness work
- Legal frameworks for AI explainability

## **3. Methodology (3-4 pages)**

- Experimental design (conditions, prompts)
- Dataset description
- Metrics definitions with rubrics
- Annotation procedures and reliability

## **4. Results (4-5 pages)**

- Quantitative findings (accuracy, interpretability metrics)
- Statistical analyses
- Visualizations

## **5. Qualitative Analysis (2-3 pages)**

- Failure case studies
- Error pattern analysis
- Representative examples

## **6. Legal and Regulatory Implications (2-3 pages)**

- Connections to legal frameworks
- Implications for accountability
- Policy recommendations

## **7. Discussion (2 pages)**

- Interpretation of findings
- Limitations
- Future work

## **8. Conclusion (1 page)**

Total length: ~20-25 pages (double-spaced)

Include appendices:

- Complete annotation rubrics
- Prompt templates
- Additional visualizations

- Inter-rater reliability calculations

**Deliverable:** Complete research paper draft

---

### Step 6.3: Prepare Presentation

**Tasks:**

Create slide deck (15-20 slides, 15-minute presentation):

1. **Title & Motivation** (1 slide)
2. **Research Question** (1 slide)
3. **Why This Matters for Law** (2 slides)
4. **Methodology Overview** (2 slides)
5. **Experimental Design** (2 slides)
6. **Key Results - Accuracy** (2 slides)
7. **Key Results - Interpretability** (2 slides)
8. **Qualitative Analysis** (2-3 slides with examples)
9. **Legal Implications** (2 slides)
10. **Limitations & Future Work** (1 slide)

#### 11. **Conclusions** (1 slide)

- Design clear, minimal slides (avoid text walls)
- Prepare speaker notes
- Practice presentation timing (aim for 12-13 minutes to allow Q&A)
- Create backup slides for anticipated questions

**Deliverable:** Presentation slides with speaker notes

---

### Step 6.4: Documentation and Reproducibility

**Tasks:**

Clean and document code:

```
python
```

```
# Add docstrings to all functions  
# Include usage examples  
# Add requirements.txt  
# Create README.md with setup instructions
```

Organize GitHub repository (if sharing):

```
repo/  
    └── README.md (setup, usage, results summary)  
    └── requirements.txt  
    └── data/  
        └── dataset_manifest.csv  
        └── responses_sample.json (not full dataset if API-generated)  
    └── notebooks/  
        └── 01_data_collection.ipynb  
        └── 02_metric_calculation.ipynb  
        └── 03_analysis_visualization.ipynb  
    └── scripts/  
        └── generate_responses.py  
        └── calculate_metrics.py  
        └── visualize_results.py  
    └── results/  
        └── tables/  
            └── figures/  
    └── paper/  
        └── final_report.pdf
```

Write comprehensive README:

- Project description
- Installation instructions
- Usage guide
- Results summary
- Citation information

Archive final outputs:

- All data files
- All code
- Final report PDF
- Presentation slides

- Supplementary materials

**Deliverable:** Complete, documented, reproducible project package

---

## Resource Requirements Summary

### Computational Resources

- **Local machine:** Sufficient for most tasks (data processing, annotation, analysis)
- **API access:** OpenAI/Anthropic OR HuggingFace Inference API
- **Estimated API cost:** \$2-10 depending on model choice and problem count

### Human Resources

- **Primary researcher** (you): ~15-20 hours/week for 6 weeks
- **Optional co-annotators:** 2-3 volunteers for 3-5 hours total (reliability checks)

### Software/Tools

- Python 3.9+ with standard ML/data science stack
- Jupyter notebooks for exploratory analysis
- LaTeX or Microsoft Word for report writing
- Presentation software (PowerPoint, Google Slides, or Beamer)

### Datasets (All Free & Public)

- GSM8K: <https://huggingface.co/datasets/openai/gsm8k>
- MATH (backup): [https://huggingface.co/datasets/hendrycks/competition\\_math](https://huggingface.co/datasets/hendrycks/competition_math)
- PRM800K (reference): <https://github.com/openai/prm800k>

### Models (Choose One Path)

#### Option A: API-based (Recommended for simplicity)

- OpenAI GPT-4o-mini (~\$2.25 for 200 problems × 3 conditions)
- Anthropic Claude 3.5 Sonnet (~\$15 for same workload)

#### Option B: Open-weight (Free but requires more setup)

- Qwen2.5-Math-7B-Instruct (strong on math)
- Llama-3.1-8B-Instruct (general purpose)
- DeepSeek-Math-7B (specialized on math)

# Risk Mitigation

## Potential Issues & Solutions

### Issue 1: API rate limits or cost overruns

- Solution: Use open-weight models via HuggingFace
- Solution: Reduce problem count to 100 minimum viable set
- Solution: Use GPT-4o-mini instead of GPT-4

### Issue 2: Response parsing difficulties

- Solution: Implement robust regex patterns with fallbacks
- Solution: Manual review and correction of edge cases
- Solution: Simplify prompts to encourage more consistent formatting

### Issue 3: Low inter-rater reliability in annotations

- Solution: Refine annotation rubrics with more specific criteria
- Solution: Conduct training session with co-annotators using examples
- Solution: Adjudicate disagreements through discussion and consensus

### Issue 4: Insufficient differences between conditions

- Solution: This is still a valid result - document null findings
- Solution: Qualitative analysis becomes more important
- Solution: Focus on failure modes and edge cases

### Issue 5: Time constraints

- Solution: Reduce annotation scope to 30-40 problems minimum
- Solution: Simplify interpretability metrics (drop one metric if needed)
- Solution: Focus on most important aspects for legal implications

---

## Success Criteria

### Minimum Viable Project

- 100+ problems evaluated across 2+ conditions
- Accuracy comparison with statistical tests
- 30+ problems annotated for step correctness
- Qualitative analysis of 10+ failure cases

- 15-page report connecting findings to legal frameworks
- 15-minute presentation

### Target Project (Full Scope)

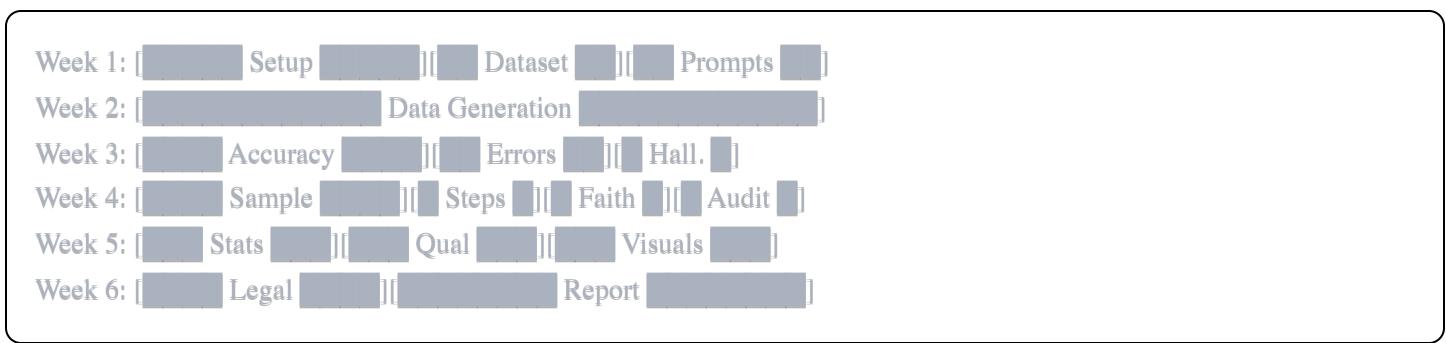
- 150-200 problems evaluated across 3 conditions
- All factual reliability metrics computed
- 50-70 problems deeply annotated
- All three interpretability metrics calculated
- Inter-rater reliability  $\geq 0.70$  (substantial agreement)
- 20-25 page comprehensive report
- Publication-quality visualizations
- Reproducible code and documentation

### Stretch Goals

- Comparison across two different models
- Analysis on both GSM8K and MATH datasets
- Integration with existing process supervision tools
- Interactive visualization dashboard
- Submission to workshop or conference

---

### Timeline Overview (Gantt Chart Format)



### Milestones:

- End of Week 2: All model responses collected ✓
- End of Week 3: Factual reliability metrics complete ✓
- End of Week 4: Interpretability annotations complete ✓
- End of Week 5: All analysis and visualizations complete ✓
- End of Week 6: Final report and presentation ready ✓

# Key Decisions to Make Now

## Decision 1: Model Selection

### Option A: OpenAI GPT-4o-mini (Recommended)

- Easy API access
- Low cost (\$2-5)
- Reliable, consistent outputs
- Closed-source (less reproducible)

### Option B: Open-weight (Qwen2.5-Math-7B)

- Free inference
- Fully reproducible
- Math-specialized
- Requires more setup
- May need GPU access

**Recommendation:** Start with GPT-4o-mini for speed, optionally add Qwen2.5-Math if time permits

---

## Decision 2: Annotation Scope

### Option A: Solo annotation (50 problems)

- Faster, no coordination overhead
- No inter-rater reliability
- Potential bias

### Option B: Multi-rater (30 problems with 2-3 raters)

- Stronger validity
- Inter-rater reliability metrics
- Requires recruiting and coordinating annotators

**Recommendation:** Start solo, recruit 1-2 volunteers for 20-30% overlap if possible

---

## Decision 3: Dataset Choice

### Option A: GSM8K only (Recommended)

- Grade-school level (easier to annotate)

- High-quality step-by-step solutions
- Standard benchmark

### Option B: MATH dataset

- More challenging (greater differentiation)
- More difficult to annotate accurately
- Requires stronger math background

**Recommendation:** Use GSM8K as primary dataset

---

### Next Immediate Actions (Today)

1. Set up Python environment and install dependencies
2. Create API account (OpenAI or Anthropic) and verify access
3. Download GSM8K dataset and explore structure
4. Sample 150-200 problems using stratified random sampling
5. Draft and test prompt templates on 5 example problems
6. Create project directory structure and initialize version control

**Time estimate:** 2-3 hours

---

### Questions to Resolve

1. Which model(s) will you use? → Decide by end of today
2. Will you have co-annotators? → Decide by Week 3
3. What is your actual weekly time budget? → Adjust timeline if needed
4. Do you have GPU access if using open models? → Determines model choice
5. What format for final submission? → Paper length, presentation format

### Conclusion

This plan provides a **realistic, achievable path** to completing your project on process vs. outcome supervision in 6 weeks. The phased approach ensures:

- ✓ Weeks 1-2: Core infrastructure and data collection complete
- ✓ Weeks 3-4: All quantitative and qualitative metrics computed

**Week 5:** Comprehensive analysis connecting technical findings

**Week 6:** Polished final deliverables with legal implications

**The plan is flexible:** You can scale down to the "Minimum Viable Project" if time is tight, or pursue "Stretch Goals" if ahead of schedule.

**Key success factors:**

- Start with reliable, easy-to-use tools (GPT-4o-mini API + GSM8K)
- Focus on quality over quantity in annotations (50 deep > 200 shallow)
- Connect every finding back to legal/regulatory implications
- Document as you go (don't wait until the end)

Good luck with your project! 